

Assignment 1

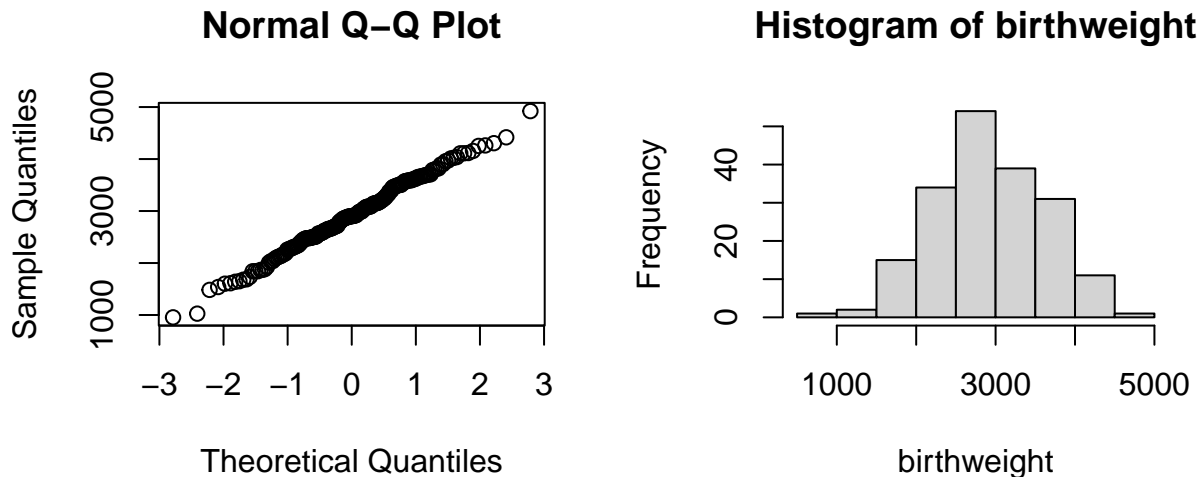
Alexia Salomons, Nathan Maxwell Jones, Yauheniya Makarevich, group 71

27 February 2023

Exercise 1.

a)

Given the QQ-plot and the histogram below, the data appears to be normal.



To calculate the CI-96% in R:

```
t.test(birthweight, conf.level = 0.96)
```

```
##
## One Sample t-test
##
## data: birthweight
## t = 57, df = 187, p-value <2e-16
## alternative hypothesis: true mean is not equal to 0
## 96 percent confidence interval:
## 2808 3019
## sample estimates:
## mean of x
## 2913
```

This gives a 96% CI of [2808-3019] with $\mu = 2913$. In order to decrease this range to 100, we can reverse the calculations in order to check the required sample size. CI is given by $\left[\bar{X} - t_{\alpha/2} \frac{s}{\sqrt{n}}, \bar{X} + t_{\alpha/2} \frac{s}{\sqrt{n}}\right]$. Thus we can conclude that $t_{\alpha/2} \frac{s}{\sqrt{n}} \leq 50$. From the data we can

determine that $s = 698$ and $t_{0.04/2} = 2.07$ (for 96% CI), and thus solve for n as follows:

```
s = sd(birthweight) # 698
t_02 = qt(1-0.02,df=188-1) # 2.07
(n = (t_02 * s / 50)^2)
```

```
## [1] 832
```

Therefore, a sample size of approximately 832 babies would be needed in order to have a 96% CI with a range of 100. We can also use the bootstrap method to calculate the 96% CI as follows:

```
B <- 1000; alpha <- 0.04; T_star <- numeric(B)

for(i in 1:B) {
  X_star <- sample(birthweight, replace = TRUE)
  T_star[i] <- mean(X_star)
}

T_star_q2 <- quantile(T_star, alpha/2)
T_star_q98 <- quantile(T_star, 1 - alpha/2)

c(2*birthweight_mean - T_star_q98, 2*birthweight_mean - T_star_q2)

## 98% 2%
## 2811 3014
```

The CI in this case is very similar to the previous one.

b) For the experts claim the hypotheses can be formulated as follows: $H_0 : \mu \leq 2800$, $H_1 : \mu > 2800$

```
t.test(birthweight, alternative = "greater", mu=2800)
```

```
##
## One Sample t-test
##
## data: birthweight
## t = 2, df = 187, p-value = 0.01
## alternative hypothesis: true mean is greater than 2800
## 95 percent confidence interval:
## 2829 Inf
## sample estimates:
## mean of x
## 2913
```

As p is smaller than 0.05, we reject the null hypothesis, supporting the claim made by the expert. The CI is infinite on the right side, since the test is one-sided, meaning that there is a 95% chance that the estimated mean is greater than 2829. For a sign test, we can use:

```
greater_weight <- as.integer(birthweight > 2800)
binom.test(sum(greater_weight), length(greater_weight), p=0.5, alt="g")
```

```
##
## Exact binomial test
```

```
##
## data:  sum(greater_weight) and length(greater_weight)
## number of successes = 107, number of trials = 188, p-value = 0.03
## alternative hypothesis: true probability of success is greater than 0.5
## 95 percent confidence interval:
##  0.507 1.000
## sample estimates:
## probability of success
##                0.569
```

We again reject H_0 , accepting H_1 , that median of the sample is bigger than 2800. Since the data appears symmetric, we assume that the mean and median are the same.

c) To compute the powers of both tests, we can assume that the data we have was sampled from a normal distribution with the same variance and $\mu = 2900$. We can then compute the probability of each test correctly rejecting $H_0 : \mu \leq 2800$ as follows. To better approximate the current situation, we will choose a sample size of $n = 188$ newborn babies.

```
B=1000; n=188 ; mu = 2900; stdev = sd(birthweight);
psign=numeric(B)
pttest=numeric(B)
for(i in 1:B) {
  x=rnorm(n, mean=mu, sd=stdev)
  pttest[i]=t.test(x, alternative = "g", mu=2800)[[3]]
  psign[i]=binom.test(sum(x>2800), n, p=0.5)[[3]]
}
```

T-test and sign test powers:

```
c(sum(pttest<0.05)/B, sum(psign<0.05)/B)
```

```
## [1] 0.634 0.373
```

We see that the power of the t-test is higher than that of the sign test. This makes sense because the sign test discards valuable information when considering only the signs, while the t-test is designed to test normal distributions, which is the case here.

d)

e) The hypotheses can be formulated as follows: $H_0 : P_{male} - P_{female} = 0$, $H_1 : P_{male} - P_{female} \neq 0$. We can use a proportion test to verify this claim.

```
prop.test(c(61, 65), c(95, 93))
```

```
##
## 2-sample test for equality of proportions with continuity correction
##
## data:  c(61, 65) out of c(95, 93)
## X-squared = 0.5, df = 1, p-value = 0.5
## alternative hypothesis: two.sided
## 95 percent confidence interval:
## -0.2016 0.0879
```


```
## sample estimates:
## prop 1 prop 2
## 0.642 0.699
```

As the p-value is bigger than 0.05, we fail to reject the null-hypothesis, meaning that there is no true difference between the mean weight of male and female babies.


Exercise 2

a) The data for *Before*, *After8weeks* and the difference between them all appear fairly normal.


Q-Q plot for Before




Histogram for Before




Q-Q plot for After




Q-Q plot for After




Q-Q plot for Difference



Q-Q plot for Difference



Plotting the cholesterol levels before and after 8 weeks for every individual yields the following:



There is a general trend that the cholesterol levels tend to decrease after 8 weeks. Individual 10 decreases less than the others. We can test whether *Before* and *After8weeks* are correlated as follows:

```
cor.test(df[, 1], df[, 2], method="pearson")
```

```
##
## Pearson's product-moment correlation
##
## data: df[, 1] and df[, 2]
## t = 29, df = 16, p-value = 2e-15
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.975 0.997
## sample estimates:
## cor
## 0.991
```

From the result of Pearson's test we can conclude that data before and data after are strongly correlated.

b) The data is paired since it is two different measurements on the same person, which also explains why they are correlated.

In order to investigate whether diet has an effect, we are going to apply two tests (Paired t-test and Permutation test) that check the difference between the means of different samples. Permutation test is applicable because it doesn't take normality into account and we are only testing for a difference between means, not how they relate to each other. First we apply the t-test as follows:

```
t.test(df[, 1], df[, 2], paired=TRUE)
```

```
##
## Paired t-test
##
## data: df[, 1] and df[, 2]
## t = 15, df = 17, p-value = 3e-11
## alternative hypothesis: true mean difference is not equal to 0
## 95 percent confidence interval:
```

```
## 0.540 0.718
## sample estimates:
## mean difference
## 0.629
```

Based on p-value obtained for paired T-test we reject null-hypothesis that samples have the same mean. Hence, there is a difference between these two samples. Next we apply the permutation test as follows:

```
diff_mean <- function(x, y) {
  return(mean(x-y))
}

stats <- diff_mean(df[, 1], df[, 2])

B <- 1000; t_star <- numeric(B)

for (i in 1:B) {
  diff_star <- t(apply(cbind(df[, 1], df[, 2]), 1, sample))
  t_star[i] <- diff_mean(diff_star[, 1], diff_star[, 2])
}

pl <- sum(t_star < stats) / B
pr <- sum(t_star > stats) / B

(p <- 2*min(pl, pr))
```

```
## [1] 0
```

This gives $p = 0$, and thus we reject the null hypothesis, again concluding that the diet has a significant effect.

c) We have sample $x_i \in Unif[3, \theta]$, $i = \overline{1, 18}$, $\theta > 3$. Analytic mean for this distribution will be $\mu = \frac{3+\theta}{2}$.

Using our sample, we can estimate the mean to be $\hat{\mu} = 5.78$ and the standard deviation to be $s = 1.1$.

Since we know that $\hat{\theta} = 2 * \hat{\mu} - 3$, we can now solve for $\hat{\theta} = 8.556$.

Then, using Central Limit Theorem, we can obtain the confidence interval for parameter θ :

$$\left[t_{-\frac{\alpha}{2}} \frac{s}{\sqrt{n}} + \frac{\theta + 3}{2}, t_{\frac{\alpha}{2}} \frac{s}{\sqrt{n}} + \frac{\theta + 3}{2} \right]$$

```
alpha <- 0.05

n <- length(df[, 2])
t_alpha <- qt(1 - alpha/2, df=n)
theta_l <- theta_hat - t_alpha*s/sqrt(n)
theta_r <- theta_hat + t_alpha*s/sqrt(n)
```

```
c(theta_l, theta_r)
```

```
## [1] 8.01 9.10
```

We can improve the CI by having more individuals in the samples, as increasing in parameter n causes decreasing in standard deviation for sample of means. Yes, we can improve the this CI because we know that we are sampling from a uniform distribution, and that we are estimating the upper bound θ . This means we can use a one-sided CI, which allows us to narrow the right side of the interval.

d) We can apply the bootstrap test as follows:

```
t <- max(df[, 2]); n <- length(df[, 2])

for (theta in 3:12) {
  B <- 1000
  t_star <- numeric(B)

  for (i in 1:B) {
    x_star <- runif(n, min = 3, max = theta)
    t_star[i] <- max(x_star)
  }

  pl <- sum(t_star < t)/B
  pr <- sum(t_star > t)/B

  p <- 2* min(pl, pr)
  print(paste("Theta =", theta, ", ", p))
}
```

```
## [1] "Theta = 3 , 0"
## [1] "Theta = 4 , 0"
## [1] "Theta = 5 , 0"
## [1] "Theta = 6 , 0"
## [1] "Theta = 7 , 0"
## [1] "Theta = 8 , 0.602"
## [1] "Theta = 9 , 0.03"
## [1] "Theta = 10 , 0"
## [1] "Theta = 11 , 0"
## [1] "Theta = 12 , 0"
```

Using bootstrap we have values of $\theta = 8$ and $\theta = 9$ for which our null hypothesis is not rejected (p-value > 0.05). We can apply Kolmogorov-Smirnov test to examine the sample since it examine if two samples were drawn from the one distribution. We can generate sample from uniform distribution and apply test on them.

```
ks.test(df[, 2], runif(100000, min = 3, max = 8))
```

```
## Warning in ks.test.default(df[, 2], runif(1e+05, min = 3, max = 8)): p-value
## will be approximate in the presence of ties
```

```
##
## Asymptotic two-sample Kolmogorov-Smirnov test
##
## data: df[, 2] and runif(1e+05, min = 3, max = 8)
## D = 0.2, p-value = 0.4
## alternative hypothesis: two-sided
```

P-value = 0.4, so we accept null hypothesis and can confirm that the sample was drawn from $U(3, \theta)$.

e) We are using sign test to test if the median of the cholesterol level is less than 6.

```
less_chol <- as.integer(df[, 2] < 6)
binom.test(sum(less_chol), length(less_chol), p=0.5, alt="g")
```

```
##
## Exact binomial test
##
## data: sum(less_chol) and length(less_chol)
## number of successes = 11, number of trials = 18, p-value = 0.2
## alternative hypothesis: true probability of success is greater than 0.5
## 95 percent confidence interval:
##  0.392 1.000
## sample estimates:
## probability of success
##                0.611
```

P-value for the test equals 0.2 which rejects alternative hypothesis. The next claim can be tested as follows:

```
less_chol <- as.integer(df[, 2] < 4.5)
binom.test(sum(less_chol), length(less_chol), p=0.25, alt="l")
```

```
##
## Exact binomial test
##
## data: sum(less_chol) and length(less_chol)
## number of successes = 3, number of trials = 18, p-value = 0.3
## alternative hypothesis: true probability of success is less than 0.25
## 95 percent confidence interval:
##  0.000 0.377
## sample estimates:
## probability of success
##                0.167
```

With a p-value of 0.3, we reject the claim that the fraction of the cholesterol levels after 8 weeks of low fat diet less than 4.5 is at most 25%.


Exercise 3

We can compute and add the variable *weight.lost* as follows:


```
df <- as.data.frame(read.table("data/diet.txt", header=TRUE))
df["weight.lost"] <- df["preweight"] - df["weight6weeks"]
head(df)
```

```
##   person gender age height preweight diet weight6weeks weight.lost
## 1      1      0  22   159        58    1        54.2         3.8
## 2      2      0  46   192        60    1        54.0         6.0
## 3      3      0  55   170        64    1        63.3         0.7
## 4      4      0  33   171        64    1        61.1         2.9
## 5      5      0  50   170        65    1        62.2         2.8
## 6      6      0  50   201        66    1        64.0         2.0
```

a) To visualize the effect of diet on weight lost, we can plot the weights before and after the diets for every individual.



We see the trend that the weights after 6 weeks of diet are lower for nearly all individuals. To test whether this trend is significant, we can apply a paired t-test as follows:

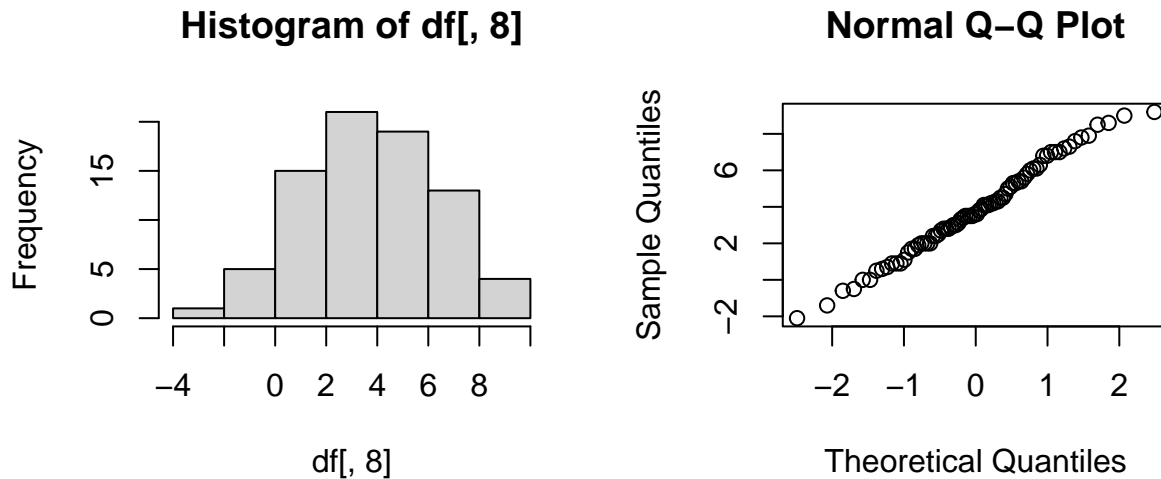
```
t.test(df[,5], df[,7], paired=TRUE)
```

```
##
## Paired t-test
##
## data: df[, 5] and df[, 7]
## t = 13, df = 77, p-value <2e-16
## alternative hypothesis: true mean difference is not equal to 0
## 95 percent confidence interval:
```

```
## 3.27 4.42
## sample estimates:
## mean difference
## 3.84
```

A p-value $< 2.2e-16$ indicates that diet does indeed have a significant effect on weight loss, with an estimated average loss of 3.84kg over the 6 week period.

To check the assumptions of the test, we need to verify that the difference between *preweight* and *weight6weeks* (ie. *weight.lost*) follows a normal distribution.



Looking at the shape of the Q-Q plot and histogram, we can conclude that *weight.lost* follows a normal distribution. Thus, the test assumptions are valid.

b) We can apply the one-way ANOVA as follows:

```
df$diet <- as.factor(df$diet)
dietaov=lm(weight.lost~diet,data=df)
anova(dietaov)
```

```
## Analysis of Variance Table
##
## Response: weight.lost
##          Df Sum Sq Mean Sq F value Pr(>F)
## diet      2     71    35.5     6.2 0.0032 **
## Residuals 75    430     5.7
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


```
summary(dietaov)
```

```
##
## Call:
## lm(formula = weight.lost ~ diet, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -5.126 -1.381 0.176 1.652 5.700
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)   3.300      0.489   6.75 2.7e-09 ***
## diet2        -0.274      0.672  -0.41  0.6845
## diet3         1.848      0.672   2.75  0.0075 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.39 on 75 degrees of freedom
## Multiple R-squared:  0.142, Adjusted R-squared:  0.119
## F-statistic:  6.2 on 2 and 75 DF,  p-value: 0.00323
```

From the ANOVA table, we can see that $p = 0.0032$, meaning that the effect of diet on weight loss is significant. From the summary table, we can see that diet 2 is worse than diet 1, however, this difference is not significant. Furthermore, we can see that diet 3 is better than diet 1 and that this difference is significant with a $p=0.0075$. Therefore, diet 3 is the best diet.

Normal Q-Q Plot



The data seems to be relatively normal, therefore, it is appropriate to use ANOVA. Generally, the Kruskal-Wallis test is used when the data does not meet the assumptions for ANOVA, even though the data meets the assumptions here, the Kruskal-Wallis can still be used:

```
kruskal.test(df$weight.lost, df$diet)

##
## Kruskal-Wallis rank sum test
##
## data: df$weight.lost and df$diet
## Kruskal-Wallis chi-squared = 10, df = 2, p-value = 0.005
```

This supports the ANOVA result as the p-value is smaller than 0.05.

c) We can apply the two-way ANOVA as follows:

```
df$gender <- as.factor(df$gender)
dietgenderaov <- lm(weight.lost~gender*diet,data=df)
```

```
anova(dietgenderaov)
```

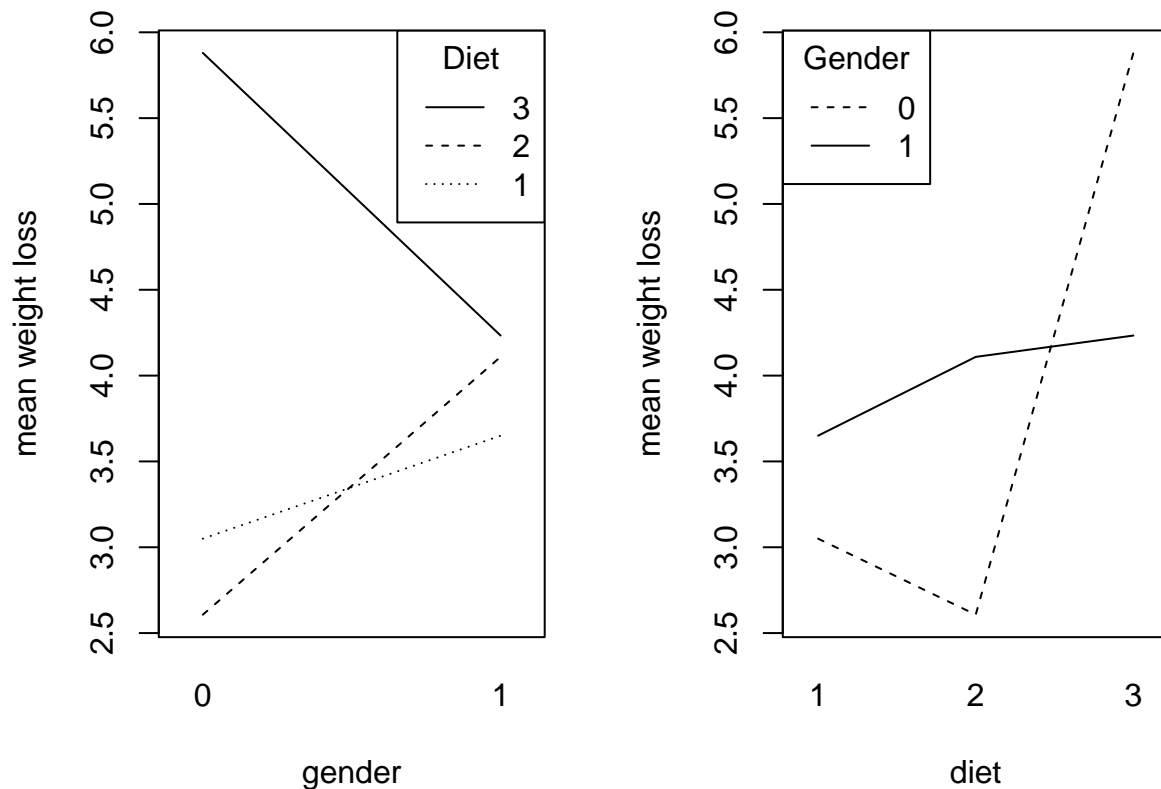
```
## Analysis of Variance Table
##
## Response: weight.lost
##           Df Sum Sq Mean Sq F value Pr(>F)
## gender      1      0    0.28    0.05 0.8206
## diet        2     60   30.21    5.62 0.0055 **
## gender:diet  2     34   16.95    3.15 0.0488 *
## Residuals   70    376    5.38
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

From this table we can see that gender does not have an effect of its own, however, it interacts with diet to affect weight loss.

```
summary(dietgenderaov)
```

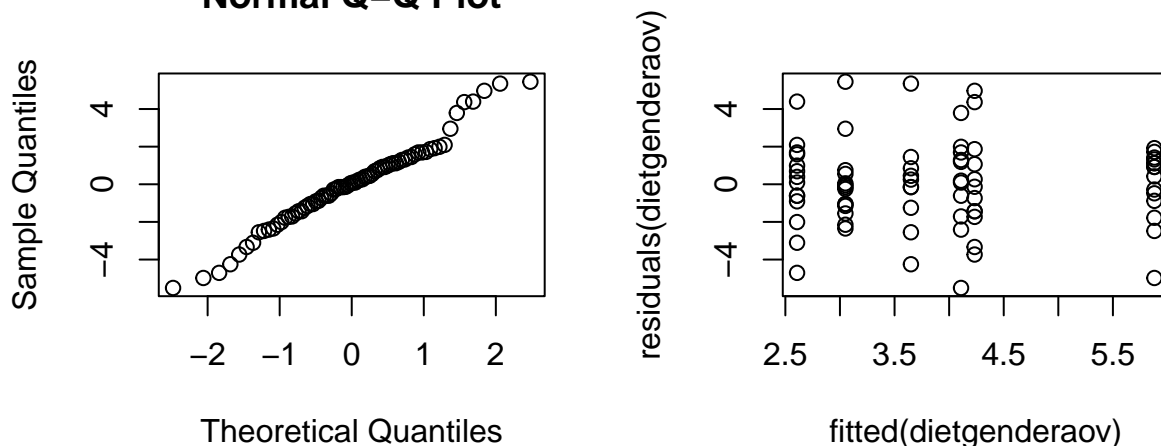
```
##
## Call:
## lm(formula = weight.lost ~ gender * diet, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.51  -1.30   0.07   1.22   5.45
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      3.050      0.620   4.92 5.5e-06 ***
## gender1          0.600      0.960   0.62  0.5340
## diet2          -0.443      0.876  -0.51  0.6149
## diet3           2.830      0.862   3.28  0.0016 **
## gender1:diet2    0.902      1.340   0.67  0.5030
## gender1:diet3   -2.247      1.315  -1.71  0.0919 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.32 on 70 degrees of freedom
## (2 observations deleted due to missingness)
## Multiple R-squared:  0.201, Adjusted R-squared:  0.144
## F-statistic: 3.52 on 5 and 70 DF, p-value: 0.00677
```

The confirms that diet 3 has the biggest effect and is significant. The interaction plots are shown below:



These graphs suggest that diet interacts with gender since the lines are not parallel. Finally, we check the model assumptions:

Normal Q-Q Plot



The residuals and QQ-plot of the gender and diet plot do not appear to be normally distributed, which violates our assumptions about normality. Therefore, one must take our previous results with caution.

e) We prefer the model from *b* as in *c* gender looks irrelevant for the weight loss and violates the needed assumptions.

```
c(dietaoav$coefficients[1],
  dietaoav$coefficients[1] + dietaoav$coefficients[2],
  dietaoav$coefficients[1] + dietaoav$coefficients[3])
```

```
## (Intercept) (Intercept) (Intercept)
##          3.30          3.03          5.15
```

An average person would lose approximately 3 kg for both diet 1 and 2, while losing approximately 5 kg for diet 3.

Exercise 4

a)

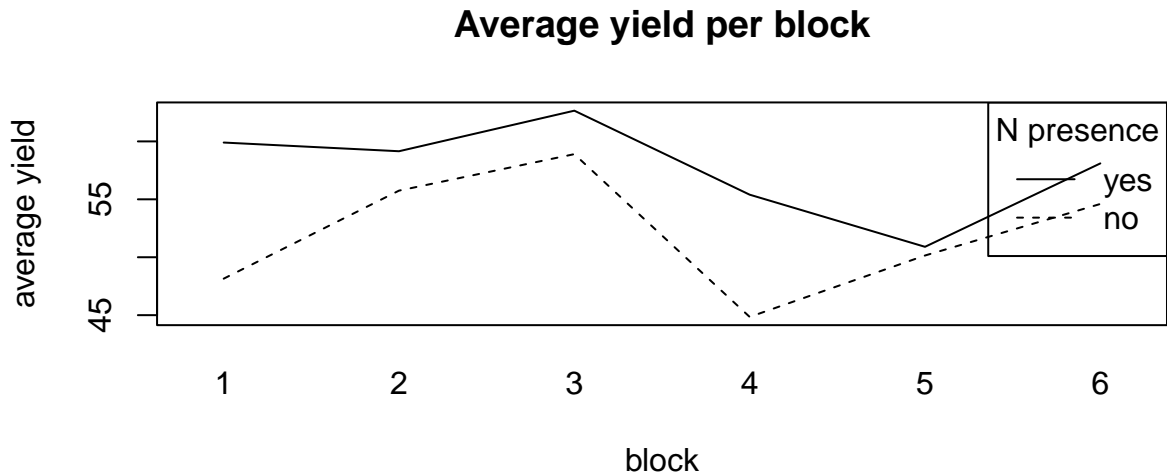
```
B <- 6; P <- 4; T <- 3
```

```
process <- c()
for (i in 1:B) {
  block <- c()
  for (tr in 1:T) {
    block <- cbind(block, as.numeric(sample(1:P) > 2))
  }
  process <- rbind(process, block)
}

process <- t(process)
rownames(process) <- c("N", "P", "K")
colnames(process) <- paste0(rep(1:6, each=4), paste0(".", rep(1:4, 6)))
process
```

```
##   1.1 1.2 1.3 1.4 2.1 2.2 2.3 2.4 3.1 3.2 3.3 3.4 4.1 4.2 4.3 4.4 5.1 5.2 5.3
## N   1   0   0   1   0   1   1   0   1   0   0   1   0   1   1   0   1   0   0
## P   0   1   0   1   0   1   0   1   0   1   0   1   1   0   1   0   1   1   0
## K   0   1   0   1   1   0   0   1   0   0   1   1   0   1   1   0   1   1   0
##   5.4 6.1 6.2 6.3 6.4
## N   1   1   0   1   0
## P   0   1   0   1   0
## K   0   0   1   1   0
```

In the table rows represent every soil additive and columns represent 6 blocks, each with 4 plots (first number - block, second number - plot). As you can see, every additive appears twice in each block.



b)

We have reason to believe that *block* may affect *yield*. This could happen because of slightly different environmental conditions: sun exposure, soil composition, etc. The plot supports this idea since it appears that average yield varies depending on the block when *N* is both present and absent.

c)

```
n_block_lm <- lm(yield~block*N, data=npk)
anova(n_block_lm)
```

```
## Analysis of Variance Table
##
## Response: yield
##           Df Sum Sq Mean Sq F value Pr(>F)
## block      5   343    68.7    3.36  0.04 *
## N          1   189   189.3    9.26  0.01 *
## block:N     5    99    19.7    0.96  0.48
## Residuals 12   245    20.4
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

It is seen from the results that *N* and *block* are significant, but their interaction is not. Let's check summary for every level to see how each block individually affect the *yield*.

```
summary(n_block_lm)
```

```
##
## Call:
## lm(formula = yield ~ block * N, data = npk)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -6.85  -1.35   0.00   1.35   6.85
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    48.15      3.20    15.06 3.7e-09 ***
## block2         7.60      4.52     1.68  0.119
```

```
## block3      10.75      4.52      2.38      0.035 *
## block4      -3.30      4.52     -0.73      0.479
## block5       2.00      4.52      0.44      0.666
## block6       6.45      4.52      1.43      0.179
## N1          11.75      4.52      2.60      0.023 *
## block2:N1    -8.35      6.39     -1.31      0.216
## block3:N1    -8.00      6.39     -1.25      0.235
## block4:N1    -1.20      6.39     -0.19      0.854
## block5:N1   -11.00      6.39     -1.72      0.111
## block6:N1    -8.25      6.39     -1.29      0.221
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.52 on 12 degrees of freedom
## Multiple R-squared:  0.72,    Adjusted R-squared:  0.464
## F-statistic: 2.81 on 11 and 12 DF,  p-value: 0.0449
```

We can see that the difference between blocks is not significant except for block 3. Additionally, it is seen that presence of *N* in the soil makes significant difference to the *yield*. Moreover, interaction between *block* and *N* is insignificant.


Because interaction is not significant in our case but *block* is significant by itself, we should go for the additive model.

```
lm <- lm(yield~block+N, data=npk)
anova(lm)
```

```
## Analysis of Variance Table
##
## Response: yield
##           Df Sum Sq Mean Sq F value Pr(>F)
## block      5    343    68.7    3.40 0.0262 *
## N          1    189   189.3    9.36 0.0071 **
## Residuals 17    344    20.2
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Using the additive model we see that both *block* and *N* are still significant. Finally, we are going to check the model assumptions about normality of residuals.

Normal Q-Q Plot



The

residuals and the Q-Q plot appear fairly normal.

Lastly, we cannot use Friedman test as we have two observations for each combination of soil additives.

d) We can investigate the pairwise interactions of N , P and K with *block* as follows:

```
npklm1 <- lm(yield~P + K + block*N, data=npk)
npklm2 <- lm(yield~N + K + block*P, data=npk)
npklm3 <- lm(yield~N + P + block*K, data=npk)
npklm4 <- lm(yield~block + N + P + K, data=npk)
```

Model 1: $Y \sim P + K + \text{block} * N$

```
anova(npklm1)
```

```
## Analysis of Variance Table
##
## Response: yield
##           Df Sum Sq Mean Sq F value Pr(>F)
## P           1      8      8.4    0.59 0.4590
## K           1     95     95.2    6.72 0.0268 *
## block        5    343     68.7    4.85 0.0164 *
## N           1    189    189.3   13.36 0.0044 **
## block:N      5     99     19.7    1.39 0.3066
## Residuals  10    142     14.2
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Model 2: $Y \sim N + K + \text{block} * P$

```
anova(npklm2)
```

```
## Analysis of Variance Table
##
## Response: yield
##           Df Sum Sq Mean Sq F value Pr(>F)
## N           1    189    189.3   11.21 0.0074 **
```

```
## K          1      95      95.2      5.64 0.0389 *
## block      5     343      68.7      4.07 0.0282 *
## P          1       8       8.4      0.50 0.4966
## block:P     5      71      14.3      0.85 0.5473
## Residuals 10     169      16.9
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Model 3: $Y \sim N + P + \text{block} * K$

```
anova(npk1m3)
```

```
## Analysis of Variance Table
##
## Response: yield
##           Df Sum Sq Mean Sq F value Pr(>F)
## N           1    189    189.3    11.14 0.0075 **
## P           1      8      8.4      0.49 0.4980
## block        5    343     68.7      4.04 0.0288 *
## K           1     95     95.2      5.60 0.0395 *
## block:K      5     70     14.1      0.83 0.5583
## Residuals   10    170     17.0
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


Model 4: $Y \sim \text{block} + N + P + K$

```
anova(npk1m4)
```

```
## Analysis of Variance Table
##
## Response: yield
##           Df Sum Sq Mean Sq F value Pr(>F)
## block        5    343     68.7      4.29 0.0127 *
## N           1    189    189.3    11.82 0.0037 **
## P           1      8      8.4      0.52 0.4800
## K           1     95     95.2      5.95 0.0277 *
## Residuals   15    240     16.0
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

We have tested interaction models as well as purely additive. All the possible interactions between *block* and soil additives are insignificant. Therefore our preference goes to the purely additive model (Model 4) as it shows the significance of independent factors with no interactions, revealing that they are all significant except *P*. Finally, we can check this model's assumptions:

Normal Q–Q Plot



e) Treating *block* as a random effect, we get the following:

```
require(lme4)
```

```
## Loading required package: lme4
```

```
## Loading required package: Matrix
```

```
npklmer <- lmer(yield~N+(1|block), REML=FALSE, data=npk)
```

```
npklmer1 <- lmer(yield~(1|block), REML=FALSE, data=npk)
```

```
anova(npklmer1, npklmer)
```

```
## Data: npk
```

```
## Models:
```

```
## npklmer1: yield ~ (1 | block)
```

```
## npklmer: yield ~ N + (1 | block)
```

```
##          npar AIC BIC logLik deviance Chisq Df Pr(>Chisq)
```

```
## npklmer1    3 159 163  -76.7      153
```

```
## npklmer     4 154 158  -72.7      146   7.9  1    0.005 **
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

From the additional analysis can be seen that *N* has a significant effect on the *yield*, which supports what we found in (c).